

Part 1 - Deployment

Learning about data pipelines

This session will be all about the components of data pipelines. Specifically, we will be using the following components for our pipeline today:

- Extractor: Filebeat, Metricbeat, and Packetbeat
- Buffer: Apache Kafka
- Transformer/Loader: Apache NiFi
- Database: OpenSearch

However, to demonstrate the pipeline working, we need an environment to gather data from. So, we will be re-deploying the 5G core network and MonArch monitoring system from yesterday.

Re-deploying the 5G Core

To begin, we will use `git clone` to fetch the source code of our data pipeline and recreate the entire Kubernetes setup we had deployed yesterday.

```
cd ~  
git clone https://github.com/hautonjt/data-pipeline  
cd data-pipeline  
./day1.sh
```

We will be extracting logs from every Kubernetes container using Filebeat, memory and CPU usage of the containers and hosts using Metricbeat, and network connection information using Packetbeat. We will also demonstrate how to extend Metricbeat to extract metrics from the Prometheus exporters in Monarch as well.

How do we create a data pipeline?

Generally, in production we want to deploy a data pipeline in the following order:

1. Database
2. Buffer
3. Transformer/Loader
4. Extractors

The database and buffer are deployed first as they only receive events and do not send them. Once both the database and buffer have been set up, connecting them with a transformer/loader is trivial. Extractors are started last to prevent a build-up of events at the buffer while waiting for the transformer/loader to initialize.

How do we create a data pipeline? (2)

However, for testing, it is easier to deploy the pipeline like this:

1. Database
2. Buffer
3. Extractors
4. Transformers/Loaders

By deploying the pipeline in this order, the extractors will send a bunch of events into the buffer first, which will allow you to inspect the structure of each event. This makes it much easier to check the events' schema as well as verify whether the information needed by the transformer is available. We will deploy our pipeline using this order.

Deploying the database

First, let's deploy the OpenSearch cluster. Run the following:

```
cd ~/data-pipeline  
./deploy-opensearch.sh
```

Note: you will need to type your password (`user`) for this script as it uses `sudo` internally

This script does the following:

- Create a highly available OpenSearch cluster of 3 pods
- Provision an instance of OpenSearch dashboards
- Configure certificate-based authentication

Switching namespaces

Both OpenSearch and OpenSearch dashboards should be deployed in a new namespace named `datapipeline`. This namespace is where all our new pipeline containers will be deployed to.

Since we will be interacting with the `datapipeline` namespace the most in this lab, we can switch the namespace that `kubectl` is connected to by running:

```
kubectl config set-context --current --namespace=datapipeline
```

Now, if you run a `kubectl` command, such as, `kubectl get pods`, it will run that command in the context of the `datapipeline` namespace by default.

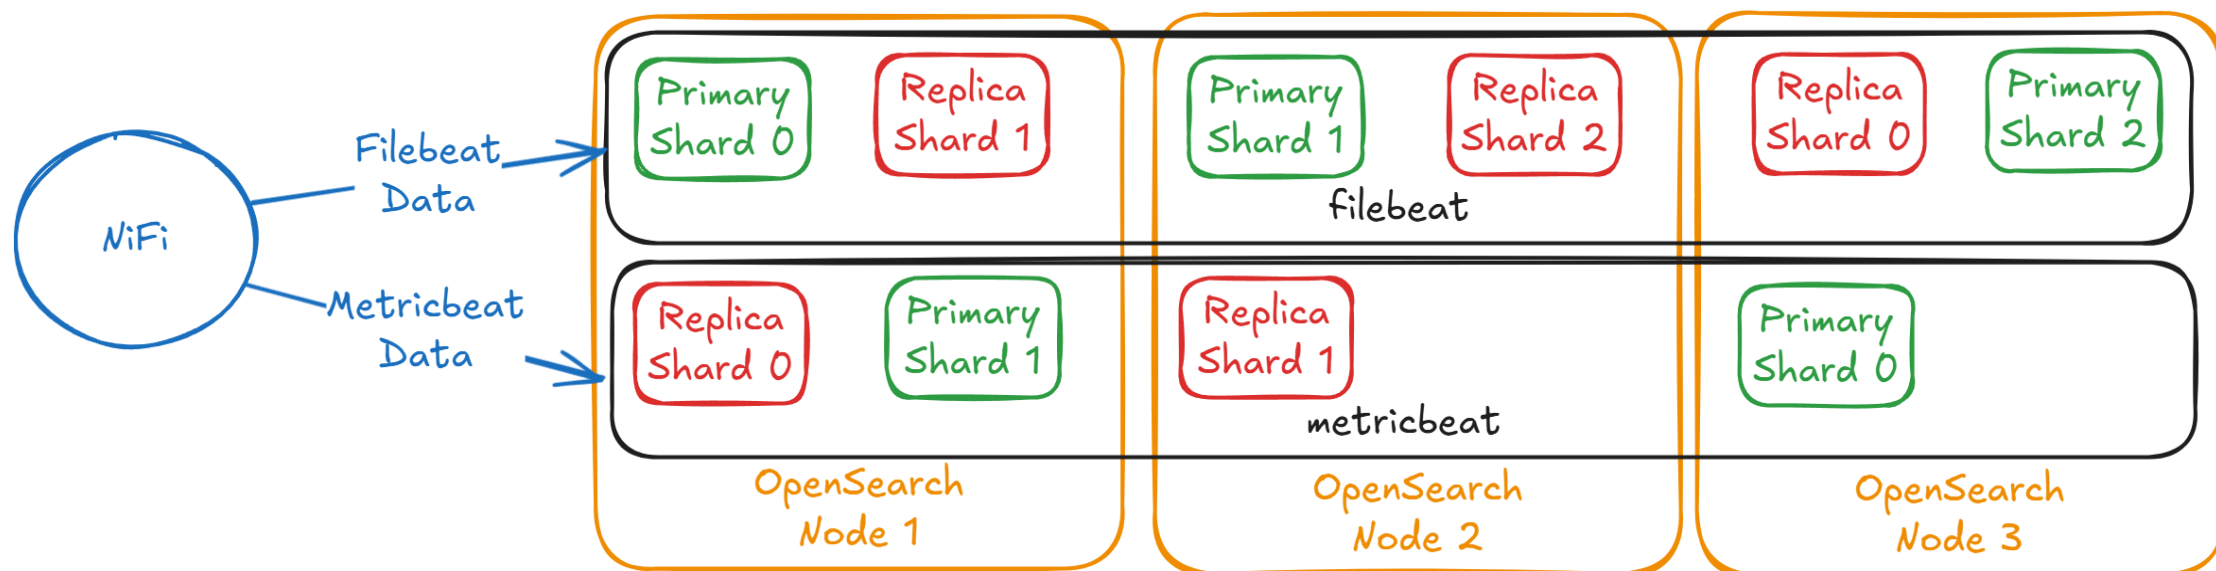
What is OpenSearch?

OpenSearch is a **fork** of Elasticsearch developed by Amazon. It was forked because the company developing Elasticsearch, Elastic NV, changed its license to restrict commercial usage.

Due to its common ancestry, many open-source tools that work with Elasticsearch also work with OpenSearch, including NiFi.

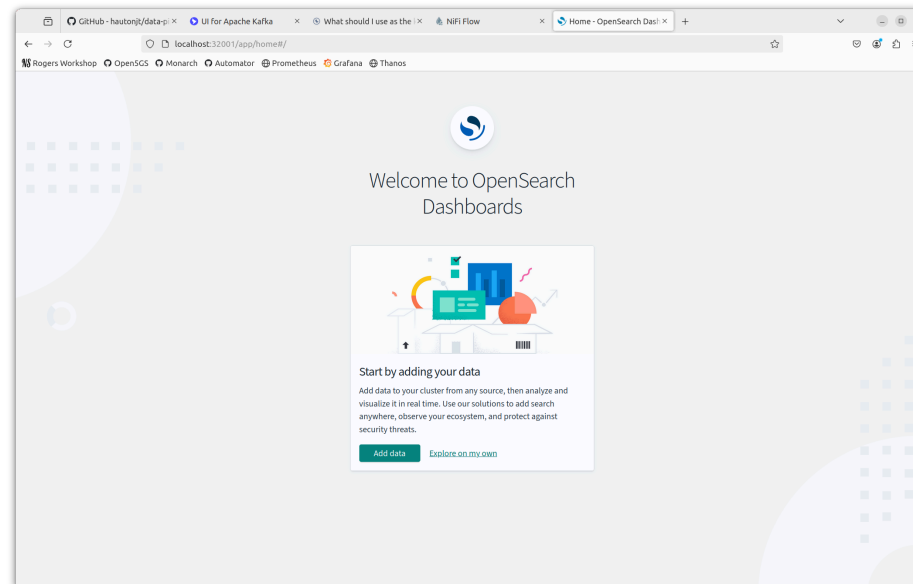
How does OpenSearch Work?

Data stored in OpenSearch using indexes. Each index consists of a number of primary shards, which allow data to be spread across multiple machines. These shards can be configured with replicas to provided high availability.



Access OpenSearch dashboards

To access the dashboards, navigate to <http://localhost:32001>. Both the username and password to log in are set to the value `admin`. If you reach the screen below, then OpenSearch has been deployed successfully. We will return to OpenSearch dashboards later on.



Deploy Kafka

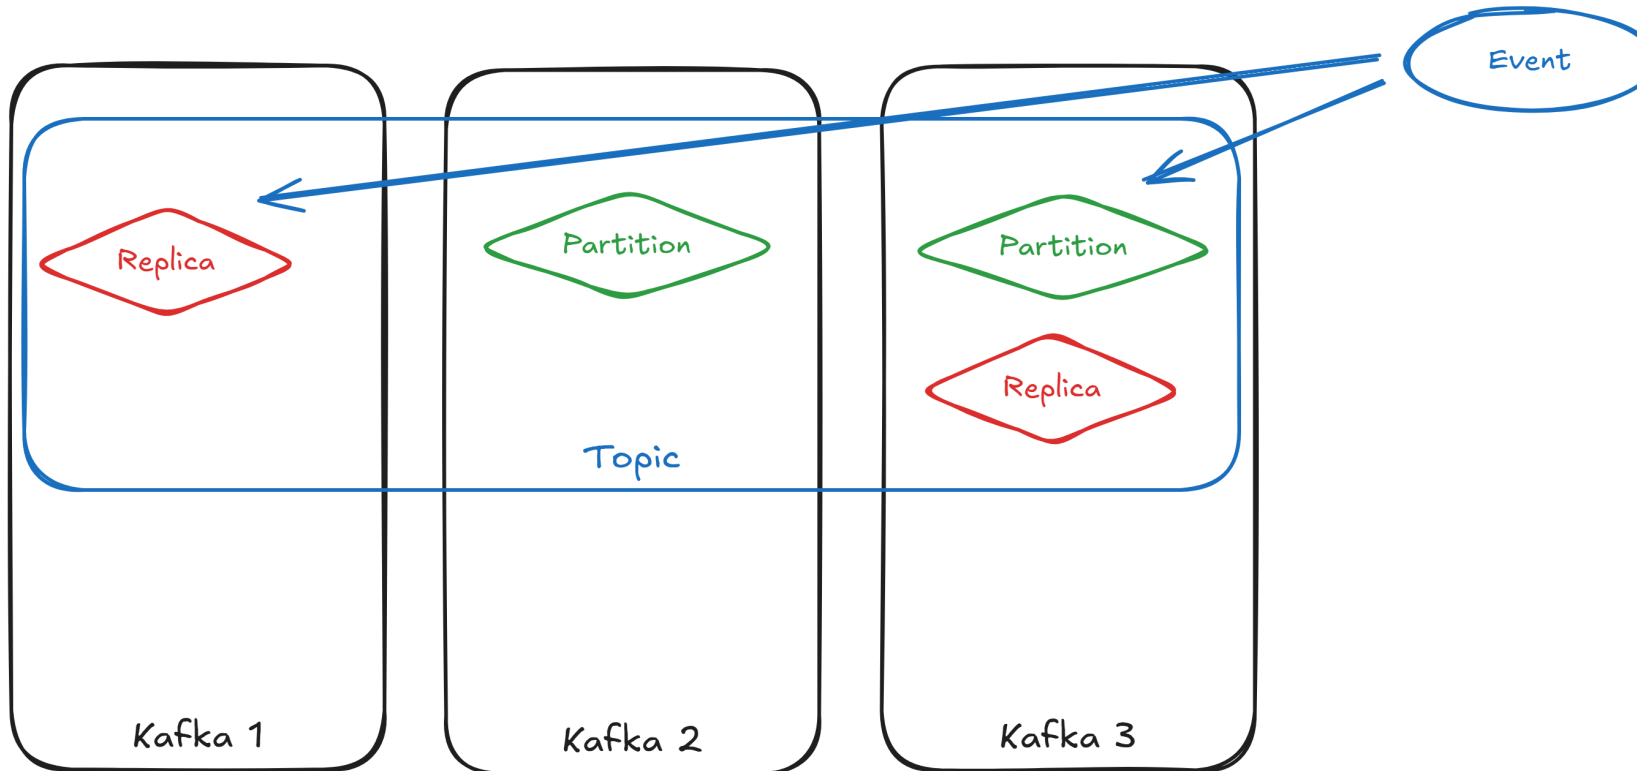
Now that OpenSearch is configured, we can install Apache Kafka. To deploy Kafka, ensure you are in the `data-pipeline` directory, and run:

```
./deploy-kafka.sh
```

This deploys both a Kafka cluster and Kafka-UI, which can be used to configure Kafka using a GUI.

What is Kafka?

Kafka is a distributed event store. Events in Kafka are mainly divided into *topics*, which are essentially categories of events. Topics are divided into partitions. These partitions allow events in a topic to be distributed across different instances of Kafka, and replicated partitions can provide fault-tolerance.



What is Kafka? (2)

Topics can be configured to have any number of partitions. The higher number of partitions, the more distributed your data becomes, at the cost of some overhead. Data redundancy and availability is configured using the replication factor, and the number of in-sync replicas.

Replication factor

Replication factor controls the number of copies of each partition. Having a replication factor of 1 means only 1 copy of each partition, meaning no redundancy. Higher replication factors mean that data from a partition can still be available as long as one copy is still present.

What is Kafka? (3)

Min in-sync replicas

Each event must be present in at least *min in-sync replicas* before being successfully written. This minimizes the probability of data loss, but also means that if there are fewer than that many replicas of a partition are present, a partition could no longer be written to.

Data persistence

Unlike message queues like RabbitMQ, data stored in Kafka topics not deleted when events are processed. Deletion is instead controlled by a data retention setting. This means that if an issue occurs in the transformer causing events to be corrupted, events can be re-ingested from the same Kafka topic before retention expires. Kafka consumers need to keep track of what events it has read from a topic.

Configuring Kafka

1. Accessing Kafka-UI

Once the deployment is done, the UI is accessible at <http://localhost:32000>.

2. Create filebeat topic

On the left sidebar, click *Topics*, then on the top right, select "Add a Topic". Fill out the following settings:

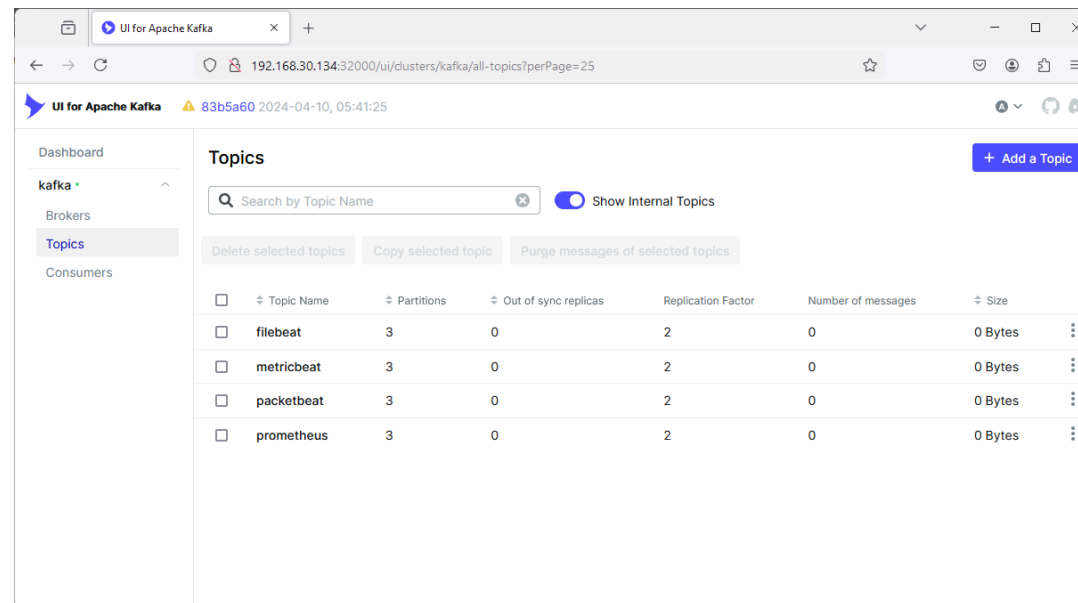
- Topic Name: filebeat
- Number of partitions: 3
- Min in-sync replicas: 2
- Replication factor: 2
- Time to retain data: 1 day (select the "1 day" box below the input)

Then select "Create".

Configuring Kafka (2)

3. Mini-Exercise

Repeat the prior steps with the topics "metricbeat", "packetbeat", and "prometheus". All other settings should remain the same. Your final screen topics screen should look something like this:



Hint: Click on the *Topics* button in the sidebar again after creating a topic to access the "Add a Topic" button quickly.

Configuring Beats Agents

Now, we will be deploying our Filebeat, Metricbeat, and Packetbeat agents. Out of the box, Beats supports Kubernetes natively, but we also want to gather metrics from the Prometheus exporters in Monarch.

To do this, we will need a configure Metricbeat to find the Prometheus endpoints.

Configuring Beats Agents (2)

In the `beats` folder in `data-pipeline`, you should be able to find a file called `metricbeat-prometheus.yaml`.

Find the section that looks like this:

```
metricbeat.modules:
  - module: prometheus
    period: 5s
    hosts: ["nssdc-kube-state-metrics.monarch.svc:8080"]
    metricsets: ["collector"]
    metrics_path: /metrics

  - module: prometheus
    period: 5s
    hosts: ["kpi-calculator-service.monarch.svc:9000"]
    metricsets: ["collector"]
    metrics_path: /metrics
```

Configuring Beats agents

These are the endpoints for Prometheus collectors. We want to add the metrics from Monarch for slice monitoring metrics.

Go to <http://localhost:30095/targets?search=>, and find the endpoints. Two of them have already been configured.

Add the AMF, SMF, and UPF endpoints on that page by replacing the `<insert AMF/SMF/UPF collector URL>` with the endpoints from Prometheus. Ensure that the port is included in the URL, but the `/metrics` subpath is removed.

Note: answers are in the `~/data-pipeline/lab/metricbeat-prometheus.yaml` file

Deploy Beats

Now we deploy Filebeat, Metricbeat, and Packetbeat all at once. To do this, run:

```
./deploy-beats.sh
```

Filebeat, Metricbeat, and Packetbeat all produce events from different types of information.

Filebeat reads informations mainly from log files and JSON HTTP endpoints. Filebeat has native integration with Kubernetes, enabling it to read the logs of all running containers, even ones deployed after Filebeat. For advanced use cases, Filebeat can also use Kubernetes annotations to separately parse the logs of certain containers if necessary.

Deploy Beats (2)

Metricbeat reads metrics, typically resource usage, from the system. It also has native integrations with various metrics providers. Notably, this includes Prometheus, allowing us to integrate with Monarch. It also integrates with kube-state-metrics, which Monarch installs.

Packetbeat captures connection information from all network interfaces present in the current running machine. This includes ingress traffic, egress traffic, and network traffic between pods.

Beats are designed to be deployed together, and have a unified schema. Thus, one advantage of using Beats is not having to worry about normalizing the schema in the transformer.

Checking Output

If you check the Kafka UI at <http://localhost:32000>, you should be able to see messages in all the topics if you navigate to Topics on the left menu. Within each individual topic, you can view the messages within the topic by selecting the "Messages" tab. If you do not see messages, please ask for help.

NiFi

Finally, we are deploying NiFi. Before deploying, To do this, run:

```
./deploy-nifi.sh
```

Note: you will need to type your password (`user`) for this script as it uses `sudo` to copy the geoip database to a system directory.

This script does the following:

- deploys a NiFi cluster of 3 nodes
- deploys a Zookeeper cluster of 3 nodes (used by NiFi for maintaining configuration and cluster information)

Warning: this deployment takes a long time to complete.

What is NiFi?

NiFi is a versatile distributed data processor. NiFi allows you to create data processing graphs that can allow you to visually see how data will be processed.

How does NiFi work?

NiFi operates using the concept of FlowFiles, which is just a container that can hold any data along with some attributes, which are essentially metadata. This means that FlowFiles inherently do not have any structure at all.

Processors act on FlowFiles and transforms them in some way. There are a vast array of processors that cater to almost any use case, and for use cases that built-in processors are incapable of handling, NiFi also supports calling external scripts as well. You can integrate ML with NiFi directly using processors instead of interfacing with Kafka.

How does NiFi work? (2)

Since FlowFiles have no structure, typically services need to be specified so the processor knows how to parse the FlowFile and what format to output it in. For example, the `JsonTreeReader` service reads the FlowFile and parses it into one or more JSON object for processing. The `JsonRecordSetWriter` writes an array of JSON objects as its output.


Services are not limited to just parsing FlowFiles, they also provide services such as SSL authentication, OpenSearch integration, caching, lookups, and more.

Part 2 - Configuration

Configuring NiFi Services

After NiFi has been deployed, we can start configuring it to consume events from Kafka and send them to OpenSearch. Once NiFi is ready, you should be able to access it by going to <http://localhost:32002/nifi>.

Before we begin, we need to define the services our processors will need. To do this, find the Settings button underneath the "NiFi Flow" text on the middle-left side of the screen. Selecting this should pop up the NiFi Flow Configuration screen. Then, select the "Controller Services" tab.



Configuring NiFi Services (2)

After that, press the "+" symbol near the top-right side of the table to add a service. We will need to add the following services:

- `JsonTreeReader`
- `JsonRecordSetWriter`
- `StandardRestrictedSSLContextService`
- `ElasticSearchClientServiceImpl`

The `JsonTreeReader` and `JsonRecordSetWriter` can both be enabled immediately by clicking on the lightning bolts to the right of their respective rows and selecting "Enable" in the window that pops up.

Note: once "Enable" has been clicked, the button will change to a "Cancel" button, so clicking again will cause the enable to be cancelled.

Configuring NiFi Services (3)

StandardRestrictedSSLContextService needs to be configured by clicking the Settings button on its row, then going to the Properties tab. Fill in the following properties:

- Keystore Filename: `keytool/keystore.p12`
- Keystore Password: `keystore`
- Key Password: (leave blank)
- Keystore Type: `PCKS12`
- Truststore Filename: `keytool/truststore.p12`
- Truststore Password: `truststore`
- Truststore Type: `PCKS12`

Then select apply.

Configuring NiFi Services (4)

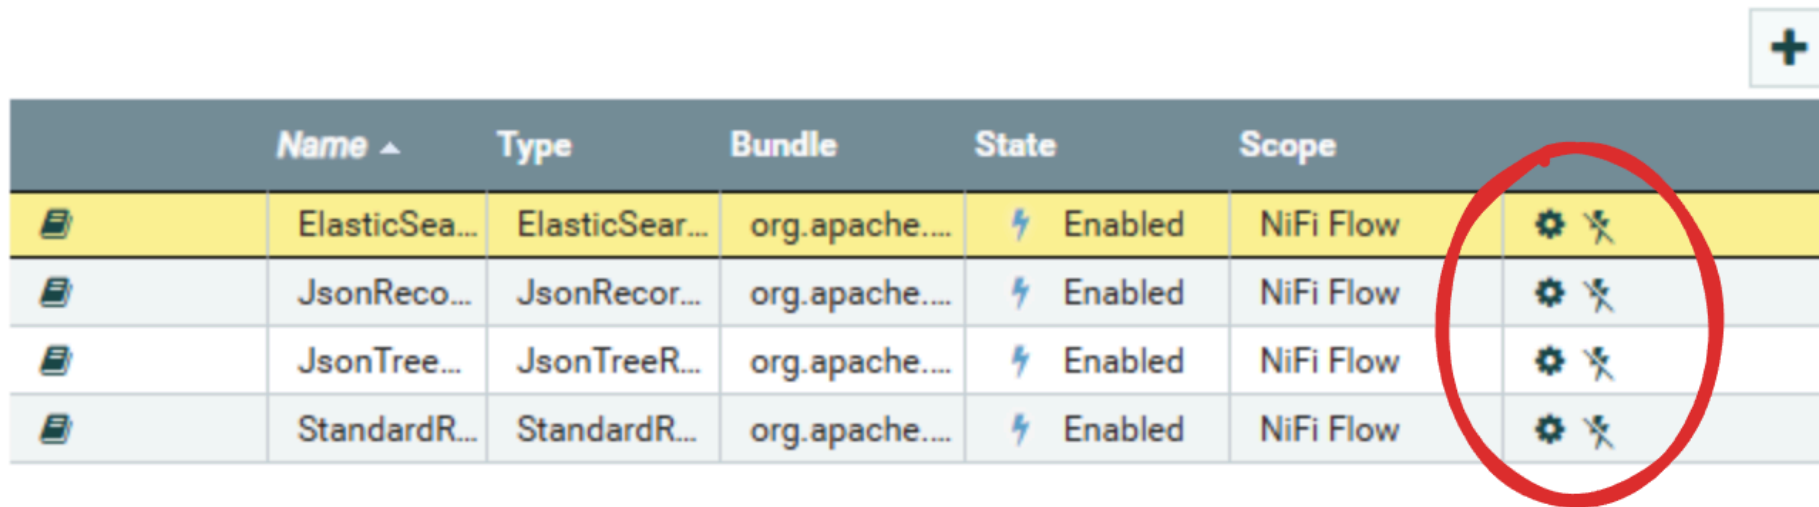
Next, we need to configure `ElasticSearchClientServiceImpl`. Click on the settings button to the right, and navigate to Properties. Enter the following settings:

















- HTTP Hosts: `https://opensearch-cluster-master:9200`
- Username: `admin`
- Password: `admin`
- SSL Context Service: `StandardRestrictedSSLContextService`

Click apply, then enable both the `StandardRestrictedSSLContextService` and the `ElasticSearchClientServiceImpl`.

Configuring NiFi Processors

Double check to ensure that all your services are properly enabled. Every service should have a crossed-out lightning bolt symbol as shown below:



	Name ▲	Type	Bundle	State	Scope	
	ElasticSea...	ElasticSear...	org.apache....	 Enabled	NiFi Flow	 
	JsonReco...	JsonRecor...	org.apache....	 Enabled	NiFi Flow	 
	JsonTree...	JsonTreeR...	org.apache....	 Enabled	NiFi Flow	 
	StandardR...	StandardR...	org.apache....	 Enabled	NiFi Flow	 

Once checked, exit out of the NiFi Flow Configuration screen. Now that all the services we need are enabled, we can now add the processors. In the top left of the screen, find the Processors icon. Click and drag the icon to create a new processor. Find the "ConsumeKafkaRecord_2_6" processor and add it.

Configuring NiFi Processors (2)

Double click on the processor to open its settings, then set the following properties:

- Kafka Brokers: `kafka:9092`
- Topic Name(s): `filebeat`
- Value Record Reader: `JsonTreeReader`
- Record Value Writer: `JsonRecordSetWriter`
- Group ID: `nifi`
- Security Protocol: `SASL_PLAINTEXT`
- SASL Mechanism: `PLAIN`
- Username: `user1`
- Password: `kafka`

Note: The username and password fields do not appear until SASL Mechanism is set to `plain`.

Configuring NiFi Processors (3)

Create a PutElasticsearchRecord processor, and configure the following settings:

- Index: `filebeat`
- Client Service: `ElasticSearchClientServiceImpl`
- Record Reader: `JsonTreeReader`

Click apply.

Configuring NiFi Relationships

NiFi processors are connected with each other using relationships. All relationships need to be configured in order for a NiFi processor to be runnable. To view a processor's relationships, double click the processor to open its settings, then navigate to the "Relationships" tab to view its relationships.

For example, the "ConsumeKafkaRecord_2_6" processor has two relationships: `success` , and `parse.failure` . Each relationship has a "terminate" and a "retry" box.

Configuring NiFi Relationships

Below each relationship is a brief explanation. For example, in our case, the event will be sent to the `parse.failure` relationship if it is not a valid JSON file. Every relationship has both a "terminate" and "retry" option.

SETTINGS	SCHEDULING	PROPERTIES	RELATIONSHIPS
----------	------------	------------	---------------

Automatically Terminate / Retry Relationships ?

`parse.failure`

☒ terminate ☐ retry

If a message from Kafka cannot be parsed using the configured Record Reader, the contents of the message will be routed to this Relationship as its own individual FlowFile.

`success`

☐ terminate ☐ retry

FlowFiles received from Kafka. Depending on demarcation strategy it is a flow file per message or a bundle of messages grouped by topic and partition.

Configuring NiFi Relationships

Enabling the "retry" option will cause the event to be re-processed by the processor a set number of times in a configurable manner. The "terminate" option causes the event to be dropped. Both "retry" and "terminate" options can be enabled, which causes the event to be dropped after all retries have been exhausted.

Any relationship that does not have the "terminate" option enabled must be connected to another processor, even if "retry" is enabled on that relationship. We will be connecting our ConsumeKafkaRecord_2_6 processor to our PutElasticsearchRecord processor.

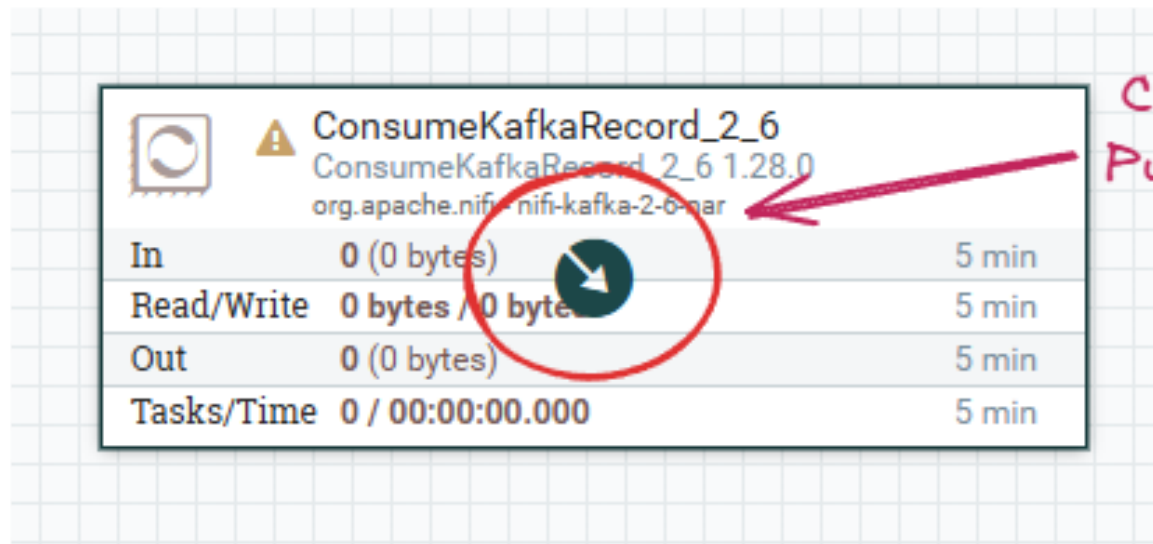
Configuring NiFi Relationships

First, we need to terminate all the relationships we don't need. Open the ConsumeKafkaRecord_2_6 settings again. Go to the "Relationships" tab, and select "terminate" under `parse.failure`, and apply.

Then, open the PutElasticsearchRecord processor settings, go to the "Relationships" tab, and select "terminate" for all relationships. We will be handling errors later. Exit out of the processor's settings.

Configuring NiFi Relationships

Now, hover over the ConsumeKafkaRecord_2_6 processor until an arrow appears in the middle of it. Drag the arrow to the PutElasticsearchRecord processor. Select the `success` relationship, and apply. This creates a connection for the `success` relationship between the ConsumeKafkaRecord_2_6 and the PutElasticsearchRecord processors.




Click and drag this arrow to the PutElasticsearchRecord processor

Running NiFi processors

The screen should now look like this:

Now, for each processor, click on it once, and in the "Operate" panel on the left, select the "Play" button to run the processor.

Note: if your processors are still displaying a yellow triangle instead of a "Stop" symbol, it has not been configured correctly. You can hover over the triangle to see the error. Please ask for help if you are stuck and the yellow triangle doesn't go away.




Running NiFi processors (2)

The screen should now look like this:

You can stop the processor again by clicking on it once, then selecting the "Stop" button in the "Operate" panel on the left.

Note: once a processor is running, its properties can no longer be edited. To edit the properties, stop the processor first. Once the processor is stopped, its properties can be edited again.



Configuring Other Processors

Duplicating processors: To duplicate a processor, right click on it and select "Copy". Then right click on any blank space and select "Paste" to paste a copy of the processor with identical configuration.

Hint: You can select multiple processors by holding down "Shift".

Mini-Exercise

Configure the processors for the *metricbeat*, *packetbeat*, and *prometheus* topics.

Ensure that both the Topic Name in the ConsumeKafkaRecord_2_6 processor and the Index in the PutElasticsearchRecord processor are set correctly.

Inspecting FlowFiles

NiFi lets you inspect queued FlowFiles pretty easily. To do this, **you must first stop the PutElasticsearchRecord processor**. Then, right click the queue (i.e., the box labeled "success"). Then select "List queue". This shows a list of all FlowFiles in the queue.

You can select the "eye" icon to view its content. **(If you get an error, ensure you stopped the PutElasticsearchRecord processor that was downstream from the queue.)**

In the "View as" selection box, select "formatted" to see the formatted JSON.

Common Issues

If no events are arriving at your ConsumeKafkaRecord_2_6 processors after a while, ensure that the settings are correct by double clicking it and verifying the processor's properties. Common mistakes include:

- Leaving the Kafka Brokers value as `localhost:9092` instead of setting it to `kafka:9092`
- Typos being present in the Topic Name(s) field

Searching Data


Now that everything is deployed, we can now visualize it in OpenSearch. Go to <http://localhost:32001>. Log in using username `admin` and password `admin` if prompted.

Select "Explore on my own". Select the menu button on the top left, scroll down and select "Dashboards Management".

Create index patterns with the names "filebeat", "packetbeat", "metricbeat", "prometheus", and "*beat,prometheus". Set the "Time field" to the value `@timestamp` for all of them.

Searching Data (2)

Open the menu, and select "Discover". Discover should look like this:



If you do not see any data, you can change the search time frame on the top right input to the left of the "Refresh" button. A relative time of "Past 1 hour" should yield some results.

You can also change the index pattern using the field on the left side of the screen, to the left of the search box.

Note: it is normal for the data in the "filebeat" index pattern to be older than the others.

Creating Visualization

Create a visualization to view data. To do this, select the menu button, and navigate to Visualize.

Create a new "Line" visualization, and choose the "prometheus-metricbeat" index pattern.

Keep the metric as count, then add an X-axis by selecting "Add" below the "Bucket" panel in the bottom right side of the screen. Add a date histogram to the X-axis using the `@timestamp` field. Press save on the top right and give it a title and optionally a description.

Feel free to some more visualizations by experimenting with other metrics.

Creating Dashboards

Create a dashboard to view multiple visualizations. Select the menu button, and navigate to Dashboards.

Select "Create new dashboard", and add the visualizations you created. Press save on the top right and give it a title and optionally a description.

NiFi Enrichment

If time permits, the NiFi enrichment slides can be found [here](#).

NiFi Error Handling

If time permits, the NiFi error handling slides can be found [here](#).

Next Steps

Congratulations!

You have successfully completed the following:

- Deployed a highly available data pipeline
- Configured agents to collect statistics from custom Prometheus exporters
- Learned about how to transform, enrich, and handle errors in data using NiFi
- Explored the data being sent by the agents using OpenSearch Dashboards

What's Next?

In the afternoon session, we will dive into [5G slice modeling and dynamic resource scaling](#).